

# Kubernetes Architecture

Ravindra Kudache

# KUBERNETES ARCHITECTURE

User Interface



kubectl

Kubernetes Master

API Server

Scheduler

Controller-Manager

etcd

Worker Node 1

Pod 1

Container 1  
Container 2  
Container 3

Pod 2

Container 1

Pod 3

Container 1  
Container 2

DOCKER

kubelet

Kube-proxy

Worker Node 2

Pod 1

Container 1  
Container 2

Pod 2

Container 1  
Container 2  
Container 3

Pod 3

Container 1

DOCKER

kubelet

Kube-proxy

# Cluster Architecture

- ❑ Kubernetes Architecture
- ❑ ETCD For Beginners
- ❑ ETCD in Kubernetes
- ❑ Kube-API Server
- ❑ Controller Managers
- ❑ Kube Scheduler
- ❑ Kubelet
- ❑ Kube Proxy



# KUBERNETES ARCHITECTURE



kube-apiserver



## Master

Manage, Plan, Schedule, Monitor  
Nodes

kubelet



## Worker Nodes

Host Application as Containers

Controller-  
Manager


D  
CLUSTER

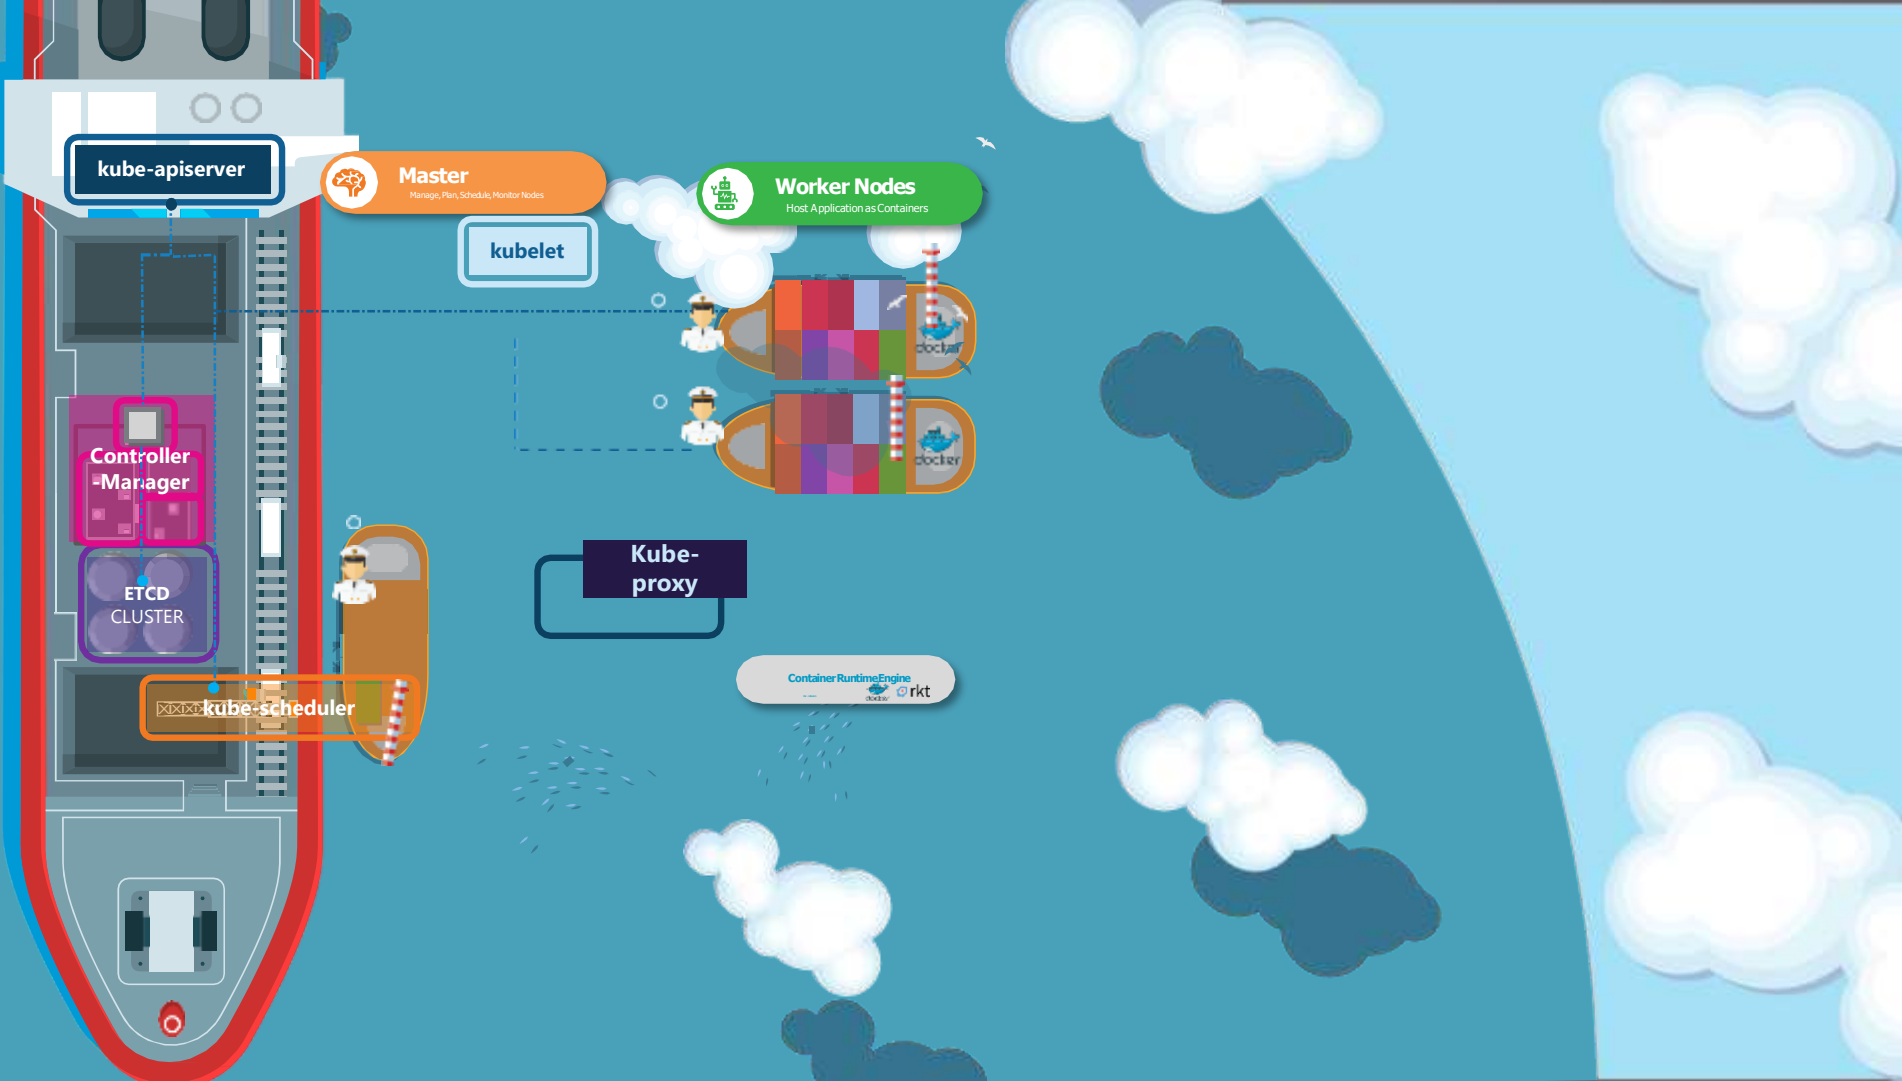
e-scheduler

Kube-proxy

Container Runtime Engine

Run containers





# Kubernetes Architecture



## Master

Manage, Plan, Schedule, Monitor  
Nodes



## Worker Nodes

Host Application as Containers

ETCD  
CLUSTER

kube-  
apiserver

Kube  
Controller  
Manager

kube-scheduler

kubelet

Kube-proxy

Container Runtime Engine

Run containers



kubelet

Kube-proxy

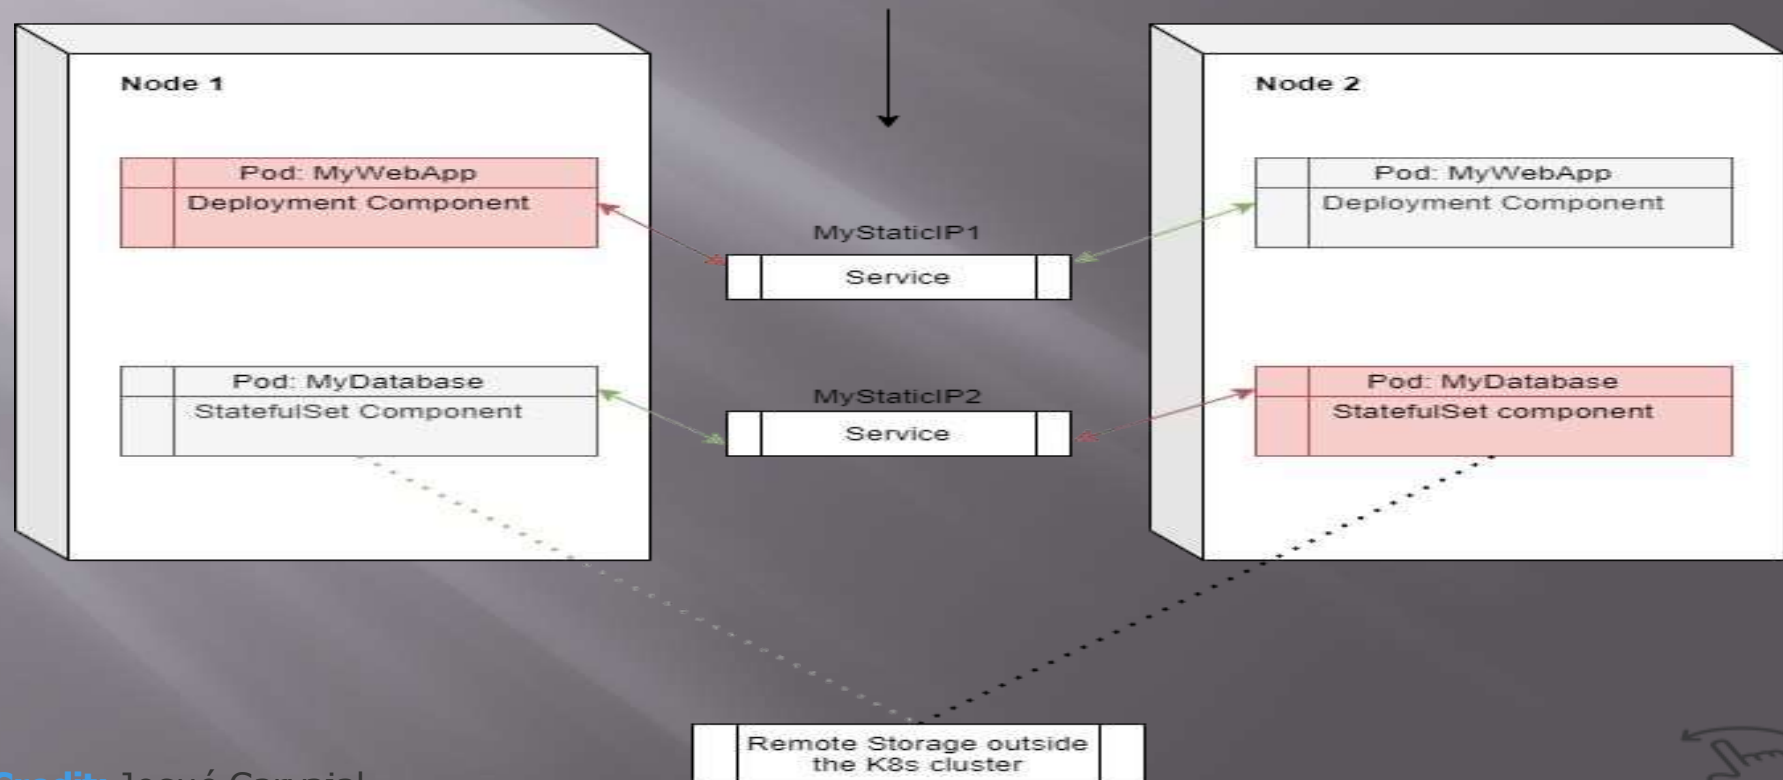
Container Runtime Engine

Run containers



# Introduction to Kubernetes Components and K8s Architecture

**New requests**  
will be distributed by the  
services to the node with  
less load





# Node

A node will contain an **N number of pods**. Usually, it is a physical server, virtual machine, or any server.



# Pod's

A Pod is the smallest unit of Kubernetes, which is an abstraction of what we knew as a container, it will run only one application on each Pod.

One important thing to point out is that **each pod has its own IP address and a new IP address will be added per creation** of a pod after it dies, this is usually something that we do not want and can be solved with Kubernetes Service.



# Service

A service allows communication between pods by adding a permanent IP to the pod, this means if the pod dies **it will conserve its IP**. It also acts as a load balancer by locating the less loaded POD and send those requests to it.



# Ingress

The ingress route the traffic into the cluster, this is usually the first service that is reached, **can be used for WEB servers that interact** with the end-user.



# Config Map

It is an external configuration file to our application that we use to configure our Pods therefore all the pods can read this configuration file **without the need of recreating or rebuilding** the pods.



## Config Map — Secrets

The secret is a configuration file that is meant to store sensitive information such as passwords. It uses **base64 encoding format**.



# Volumes

Each time a node is restarted the data on it is deleted, to avoid this Kubernetes has the Volumes component, **which is meant to get persistence.**



# Pod Blueprint

Defines the pod and its replicating mechanisms.





# Deployment

This is how we deploy new pods, it is meant for stateless applications (**applications that do not store data**).



# Stateful

- ▣ This is meant for stateful applications **such as databases**, this must be used in this way since it avoids data inconsistency when two replicated nodes sharing a database pod are accessing the same database, how do Kubernetes can ensure there is data consistency, well by using stateful component and making sure who is using the storage at a given point of time.





# ETCD

## FOR BEGINNERS

# I Objectives

- What is ETCD?
- What is a Key-Value Store?
- How to get started quickly?
- How to operate ETCD?
- What is a distributed system?
- How ETCD Operates
- RAFT Protocol
- Best practices on number of nodes

ETCD is a distributed reliable key-value store that  
is Simple, Secure & Fast

# | key-value store

Tabular/Relational Databases

Name	Age	Location	Salary	Grade
John Doe	45	New York	5000	
Dave Smith	34	New York	4000	
Aryan Kumar	10	New York		A
Lauren Rob	13	Bangalore		C
Lily Oliver	15	Bangalore		B

# key-value store

Key	Value
Name	John Doe
Age	45
Location	New York
Salary	5000

Key	Value
Name	Dave Smith
Age	34
Location	New York
Salary	4000
Organization	ACME

Key	Value
Name	Aryan Kumar
Age	10
Location	New York
Grade	A

Key	Value
Name	Lauren Rob
Age	13
Location	Bangalore
Grade	C

Key	Value
Name	Lily Oliver
Age	15
Location	Bangalore
Grade	B

# key-value store

```
{  
  "name": "John Doe",  
  "age": 45,  
  "location": "New York",  
  "salary": 5000  
}
```

```
{  
  "name": "Dave Smith",  
  "age": 34,  
  "location": "New York",  
  "salary": 4000,  
  "organization": "ACME"  
}
```

```
{  
  "name": "Aryan Kumar",  
  "age": 10,  
  "location": "New York",  
  "Grade": "A"  
}
```

```
{  
  "name": "Lily Oliver",  
  "age": 15,  
  "location": "Bangalore",  
  "Grade": "B"  
}
```

```
{  
  "name": "Lauren Rob",  
  "age": 13,  
  "location": "Bangalore",  
  "Grade": "C"  
}
```



# ETCD

In  
Kubernetes


An abstract network diagram consisting of several white dots connected by thin white lines, forming a web-like structure. The dots are positioned at various points, with some acting as central hubs connected to multiple other nodes. The lines are thin and white, contrasting with the dark background.



## Master

Manage, Plan, Schedule, Monitor  
Nodes

ETCD  
CLUSTER



# Explore ETCD

```
kubectl exec etcd-master -n kube-system etcdctl get / --prefix --keys-only  
  
/registry/apiregistration.k8s.io/apiservices/v1.  
/registry/apiregistration.k8s.io/apiservices/v1.apps  
/registry/apiregistration.k8s.io/apiservices/v1.authentication.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.authorization.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.autoscaling  
/registry/apiregistration.k8s.io/apiservices/v1.batch  
/registry/apiregistration.k8s.io/apiservices/v1.networking.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.rbac.authorization.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1.storage.k8s.io  
/registry/apiregistration.k8s.io/apiservices/v1beta1.admissionregistration.k8s.io
```

Registry

minions

pods


replicasets

deployments

roles

secrets

# ETCD in HA Environment



## etcd.service

```
ExecStart=/usr/local/bin/etcd \\  
  --name ${ETCD_NAME} \\  
  --cert-file=/etc/etcd/kubernetes.pem \\  
  --key-file=/etc/etcd/kubernetes-key.pem \\  
  --peer-cert-file=/etc/etcd/kubernetes.pem \\  
  --peer-key-file=/etc/etcd/kubernetes-key.pem \\  
  --trusted-ca-file=/etc/etcd/ca.pem \\  
  --peer-trusted-ca-file=/etc/etcd/ca.pem \\  
  --peer-client-cert-auth \\  
  --client-cert-auth \\  
  --initial-advertise-peer-urls https://${INTERNAL_IP}:2380 \\  
  --listen-peer-urls https://${INTERNAL_IP}:2380 \\  
  --listen-client-urls https://${INTERNAL_IP}:2379,https://127.0.0.1:2379 \\  
  --advertise-client-urls https://${INTERNAL_IP}:2379 \\  
  --initial-cluster-token etcd-cluster-0 \\  
  --initial-cluster controller-0=https://${CONTROLLER0_IP}:2380,controller-1=https://${CONTROLLER1_IP}:2380 \\  
  --initial-cluster-state new \\  
  --data-dir=/var/lib/etcd
```

kube-api server



kube-apiserver



## Master

Manage, Plan, Schedule, Monitor  
Nodes

kubelet




## Worker Nodes

Host Application as Containers

Controller-  
Manager

CD  
CLUSTER

e-scheduler



# Kubernetes Architecture



```
kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
master	Ready	master	20m	v1.11.3
node01	Ready	<none>	20m	v1.11.3

kube-apiserver

Controller-manager

ETCD CLUSTER

kube-scheduler



**Master**

Manage, Plan, Schedule, Monitor Nodes

kubelet



**Worker Nodes**

Host Application as Containers

Container Runtime Engine

Run containers



**Worker Nodes**

Host Application as Containers

Container Runtime Engine

Run containers



# Kubernetes Architecture



```
curl -X POST /api/v1/namespaces/default/pods ...[other]
```

Pod created!

kube-apiserver

Controller-  
manager

ETCD  
CLUSTER

kube-scheduler



**Master**

Manage, Plan, Schedule, Monitor  
Nodes



**Worker Nodes**

Host Application as Containers

kubelet

**Container Runtime Engine**

Run containers



**Worker Nodes**

Host Application as Containers

kubelet

**Container Runtime Engine**

Run containers





# Installing kube-api server

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kube-apiserver
```

## kube-apiserver.service

```
ExecStart=/usr/local/bin/kube-apiserver \\  
  --advertise-address=${INTERNAL_IP} \\  
  --allow-privileged=true \\  
  --apiserver-count=3 \\  
  --authorization-mode=Node,RBAC \\  
  --bind-address=0.0.0.0 \\  
  --enable-admission-  
plugins=Initializers,NamespaceLifecycle,NodeRestriction,LimitRanger,ServiceAccount,DefaultStorageClass,ResourceQuota \\  
  --enable-swagger-ui=true \\  
  --etcd-servers=https://127.0.0.1:2379 \\  
  --event-ttl=1h \\  
  --experimental-encryption-provider-config=/var/lib/kubernetes/encryption-config.yaml \\  
  --runtime-config=api/all \\  
  --service-account-key-file=/var/lib/kubernetes/service-account.pem \\  
  --service-cluster-ip-range=10.32.0.0/24 \\  
  --service-node-port-range=30000-32767 \\  
  --v=2
```

# View api-server - kubeadm

```
kubectl get pods -n kube-system
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcd6894-hwrq9	1/1	Running	0	16m
kube-system	coredns-78fcd6894-rzhjr	1/1	Running	0	16m
kube-system	etcd-master	1/1	Running	0	15m
kube-system	kube-apiserver-master	1/1	Running	0	15m
kube-system	kube-controller-manager-master	1/1	Running	0	15m
kube-system	kube-proxy-lzt6f	1/1	Running	0	16m
kube-system	kube-proxy-zm5qd	1/1	Running	0	16m
kube-system	kube-scheduler-master	1/1	Running	0	15m
kube-system	weave-net-29z42	2/2	Running	1	16m
kube-system	weave-net-snm1	2/2	Running	1	16m

# View api-server options - kubeadm

```
cat /etc/kubernetes/manifests/kube-apiserver.yaml
```

```
spec:
  containers:
  - command:
    - kube-apiserver
    - --authorization-mode=Node,RBAC
    - --advertise-address=172.17.0.32
    - --allow-privileged=true
    - --client-ca-file=/etc/kubernetes/pki/ca.crt
    - --disable-admission-plugins=PersistentVolumeLabel
    - --enable-admission-plugins=NodeRestriction
    - --enable-bootstrap-token-auth=true
    - --etcd-cafile=/etc/kubernetes/pki/etcd/ca.crt
    - --etcd-certfile=/etc/kubernetes/pki/apiserver-etcd-client.crt
    - --etcd-keyfile=/etc/kubernetes/pki/apiserver-etcd-client.key
    - --etcd-servers=https://127.0.0.1:2379
    - --insecure-port=0
    - --kubelet-client-certificate=/etc/kubernetes/pki/apiserver-kubelet-client.crt
    - --kubelet-client-key=/etc/kubernetes/pki/apiserver-kubelet-client.key
    - --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname
    - --proxy-client-cert-file=/etc/kubernetes/pki/front-proxy-client.crt
    - --proxy-client-key-file=/etc/kubernetes/pki/front-proxy-client.key
    - --requestheader-allowed-names=front-proxy-client
    - --requestheader-client-ca-file=/etc/kubernetes/pki/front-proxy-ca.crt
    - --requestheader-extra-headers-prefix=X-Remote-Extra-
    - --requestheader-group-headers=X-Remote-Group
    - --requestheader-username-headers=X-Remote-User
```

# View api-server options

```
cat /etc/systemd/system/kube-apiserver.service
```

```
[Service]
ExecStart=/usr/local/bin/kube-apiserver \\\
  --advertise-address=${INTERNAL_IP} \\\
  --allow-privileged=true \\\
  --apiserver-count=3 \\\
  --audit-log-maxage=30 \\\
  --audit-log-maxbackup=3 \\\
  --audit-log-maxsize=100 \\\
  --audit-log-path=/var/log/audit.log \\\
  --authorization-mode=Node,RBAC \\\
  --bind-address=0.0.0.0 \\\
  --client-ca-file=/var/lib/kubernetes/ca.pem \\\
  --enable-admission-
plugins=Initializers,NamespaceLifecycle,NodeRestriction,LimitRanger,ServiceAccount,DefaultStorageClass,ResourceQuota \\\
  --enable-swagger-ui=true \\\
  --etcd-cafile=/var/lib/kubernetes/ca.pem \\\
  --etcd-certfile=/var/lib/kubernetes/kubernetes.pem \\\
  --etcd-keyfile=/var/lib/kubernetes/kubernetes-key.pem \\\
  --etcd-
servers=https://10.240.0.10:2379,https://10.240.0.11:2379,https://10.240.0.12:2379 \\\
  --event-ttl=1h \\\
  --experimental-encryption-provider-config=/var/lib/kubernetes/encryption-config.yaml \\\
  --kubelet-certificate-authority=/var/lib/kubernetes/ca.pem \\\
  --kubelet-client-certificate=/var/lib/kubernetes/kubernetes.pem \\\
```

# View api-server options

```
ps -aux | grep kube-apiserver
```

```
root      2348   3.3 15.4 399040 315604 ?        Ssl  15:46   1:22 kube-apiserver --authorization-mode=Node,RBAC --  
advertise-address=172.17.0.32 --allow-privileged=true --client-ca-file=/etc/kubernetes/pki/ca.crt --disable-  
admission-plugins=PersistentVolumeLabel --enable-admission-plugins=NodeRestriction --enable-bootstrap-token-  
auth=true --etcd-cafile=/etc/kubernetes/pki/etcd/ca.crt --etcd-certfile=/etc/kubernetes/pki/apiserver-etcd-  
client.crt --etcd-keyfile=/etc/kubernetes/pki/apiserver-etcd-client.key --etcd-servers=https://127.0.0.1:2379 --  
insecure-port=0 --kubelet-client-certificate=/etc/kubernetes/pki/apiserver-kubelet-client.crt --kubelet-client-  
key=/etc/kubernetes/pki/apiserver-kubelet-client.key --kubelet-preferred-address-  
types=InternalIP,ExternalIP,Hostname --proxy-client-cert-file=/etc/kubernetes/pki/front-proxy-client.crt --proxy-  
client-key-file=/etc/kubernetes/pki/front-proxy-client.key --requestheader-allowed-names=front-proxy-client --  
requestheader-client-ca-file=/etc/kubernetes/pki/front-proxy-ca.crt --requestheader-extra-headers-prefix=X-Remote-  
Extra- --requestheader-group-headers=X-Remote-Group --requestheader-username-headers=X-Remote-User --secure-  
port=6443 --service-account-key-file=/etc/kubernetes/pki/sa.pub --service-cluster-ip-range=10.96.0.0/12 --tls-  
cert-file=/etc/kubernetes/pki/apiserver.crt --tls-private-key-file=/etc/kubernetes/pki/apiserver.key
```

# Kube Controller Manager



## Master

Manage, Plan, Schedule, Monitor  
Nodes

## Controller-Manager


ETCD  
CLUSTER

kube-scheduler

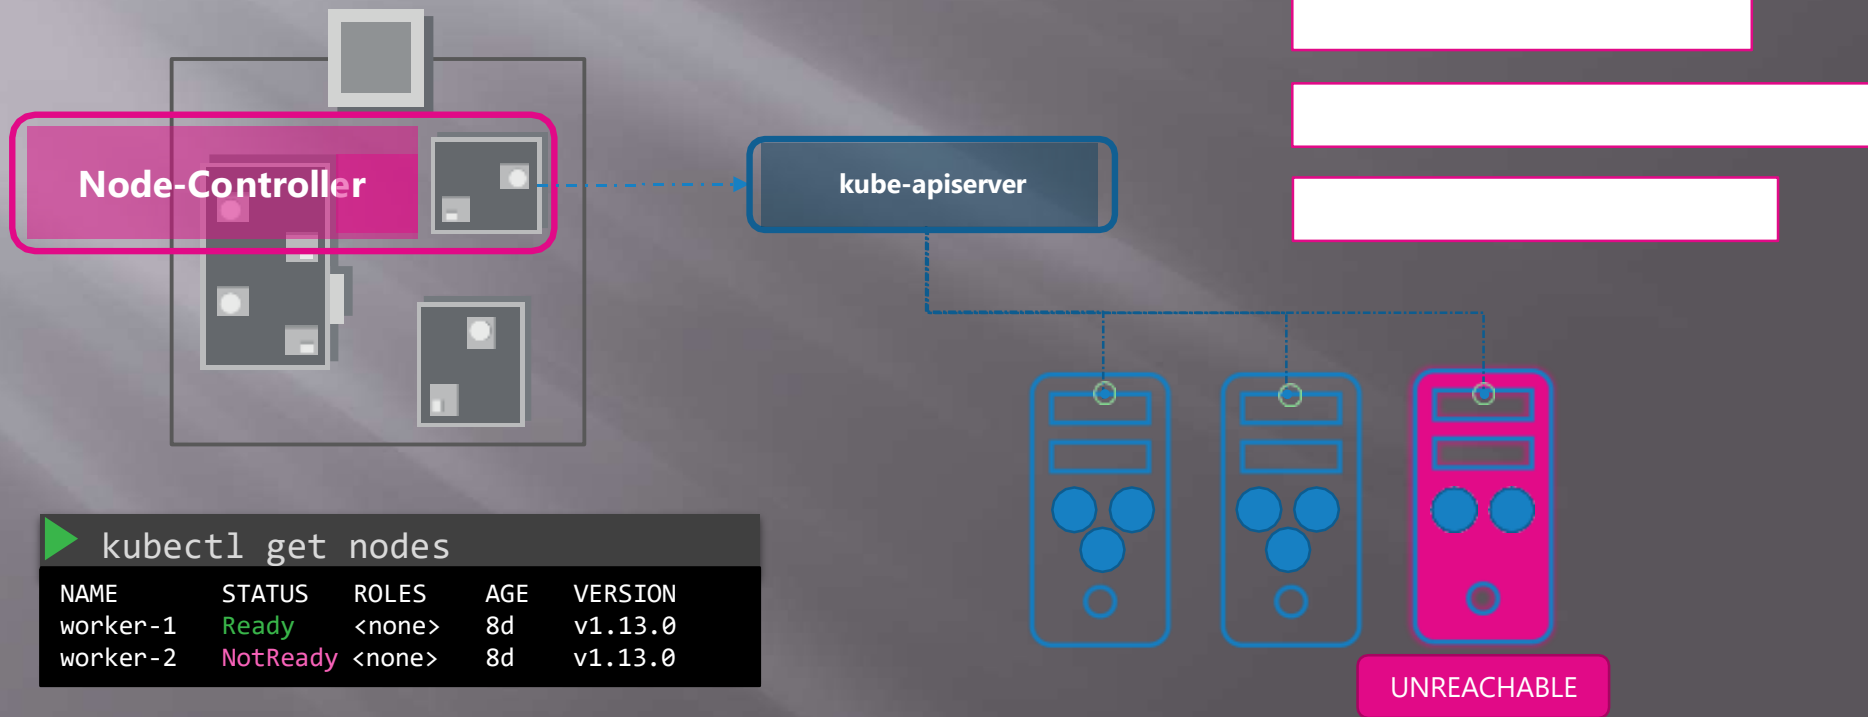


## Worker Nodes

Host Application as Containers




# Controller

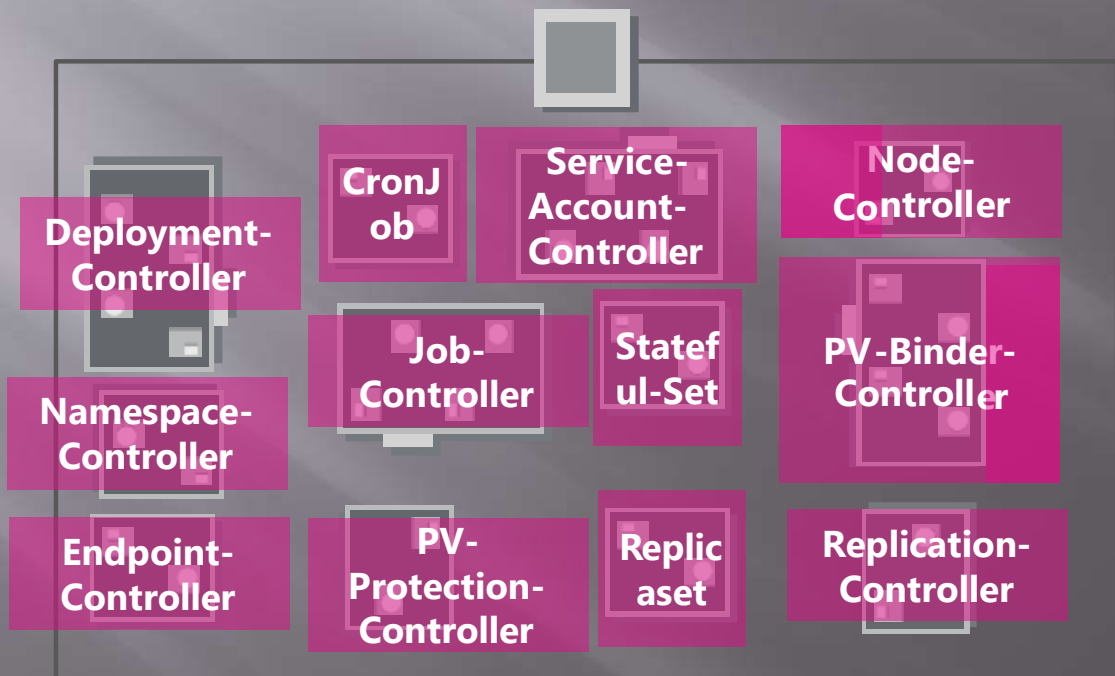




# Controller



# Controller



[Redacted]


[Redacted]

[Redacted]

[Redacted]

[Redacted]

# | Controller



# Installing kube-controller-manager

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kube-controller-manager
```

## kube-controller-manager.service

```
ExecStart=/usr/local/bin/kube-controller-manager \\  
  --address=0.0.0.0 \\  
  --cluster-cidr=10.200.0.0/16 \\  
  --cluster-name=kubernetes \\  
  --cluster-signing-cert-file=/var/lib/kubernetes/ca.pem \\  
  --cluster-signing-key-file=/var/lib/kubernetes/ca-key.pem \\  
  --kubeconfig=/var/lib/kubernetes/kube-controller-manager.kubeconfig \\  
  --leader-elect=true \\  
  --root-ca-file=/var/lib/kubernetes/ca.pem \\  
  --service-account-private-key-file=/var/lib/kubernetes/service-account-key.pem \\  
  --service-cluster-ip-range=10.32.0.0/24 \\  
  --use-service-account-credentials=true \\  
  --v=2
```

```
--node-monitor-period=5s
```

```
--controllers stringSlice      Default: [*]
```

A list of controllers to enable. '\*' enables all on-by-default controllers, 'foo' enables the controller named 'foo', '-foo' disables the controller named 'foo'.

All controllers: attachdetach, bootstrapsigner, clusterrole-aggregation, cronjob, csrapproving, csrcleaner, csrsigning, daemonset, deployment, disruption, endpoint, garbagecollector, horizontalpodautoscaling, job, namespace, nodeipam, nodelifecycle, persistentvolume-binder, persistentvolume-expander, podgc, pv-protection, pvc-protection, replicaset, replicationcontroller,

```
d=40s
```

```
s
```

# Installing kube-controller-manager

```
--controllers stringSlice      Default: [*]  
A list of controllers to enable. '*' enables all on-by-default controllers, 'foo' enables the  
controller named 'foo', '-foo' disables the controller named 'foo'.  
All controllers: attachdetach, bootstrapsigner, clusterrole-aggregation, cronjob, csrapproving,  
csrcleaner, csrsigning, daemonset, deployment, disruption, endpoint, garbagecollector,  
horizontalpodautoscaling, job, namespace, nodeipam, nodelifecycle, persistentvolume-binder,  
persistentvolume-expander, podgc, pv-protection, pvc-protection, replicaset, replicationcontroller,  
resourcequota, root-ca-cert-publisher, route, service, serviceaccount, serviceaccount-token,  
statefulset, tokencleaner, ttl, ttl-after-finished  
Disabled-by-default controllers: bootstrapsigner, tokencleaner
```

# View kube-controller-manager - kubeadm

```
kubectl get pods -n kube-system
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcd6894-hwrq9	1/1	Running	0	16m
kube-system	coredns-78fcd6894-rzhjr	1/1	Running	0	16m
kube-system	etcd-master	1/1	Running	0	15m
kube-system	kube-apiserver-master	1/1	Running	0	15m
kube-system	kube-controller-manager-master	1/1	Running	0	15m
kube-system	kube-proxy-lzt6f	1/1	Running	0	16m
kube-system	kube-proxy-zm5qd	1/1	Running	0	16m
kube-system	kube-scheduler-master	1/1	Running	0	15m
kube-system	weave-net-29z42	2/2	Running	1	16m
kube-system	weave-net-snm1	2/2	Running	1	16m

□	NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
□	kube-system	coredns-78fcd6894-hwrq9	1/1	Running	0	16m
	kube-system	etcd-master	1/1	Running	0	15m
	kube-system	kube-controller-manager-master	1/1	Running	0	15m
	kube-system	kube-proxy-lzt6f	1/1	Running	0	16m
	kube-system	kube-proxy-zm5qd	1/1	Running	0	16m
	kube-system	weave-net-29z42	2/2	Running	1	16m
	kube-system	weave-net-snm1	2/2	Running	1	16m

# View kube-controller-manager options

## - kubeadm

```
cat /etc/kubernetes/manifests/kube-controller-manager.yaml
```

```
spec:
  containers:
  - command:
    - kube-controller-manager
    - --address=127.0.0.1
    - --cluster-signing-cert-file=/etc/kubernetes/pki/ca.crt
    - --cluster-signing-key-file=/etc/kubernetes/pki/ca.key
    - --controllers=*,bootstrapsigner,tokencleaner
    - --kubeconfig=/etc/kubernetes/controller-manager.conf
    - --leader-elect=true
    - --root-ca-file=/etc/kubernetes/pki/ca.crt
    - --service-account-private-key-file=/etc/kubernetes/pki/sa.key
    - --use-service-account-credentials=true
```

# View controller-manager options

```
cat /etc/systemd/system/kube-controller-manager.service
```

```
[Service]
ExecStart=/usr/local/bin/kube-controller-manager \\\
  --address=0.0.0.0 \\\
  --cluster-cidr=10.200.0.0/16 \\\
  --cluster-name=kubernetes \\\
  --cluster-signing-cert-file=/var/lib/kubernetes/ca.pem \\\
  --cluster-signing-key-file=/var/lib/kubernetes/ca-key.pem \\\
  --kubeconfig=/var/lib/kubernetes/kube-controller-manager.kubeconfig \\\
  --leader-elect=true \\\
  --root-ca-file=/var/lib/kubernetes/ca.pem \\\
  --service-account-private-key-file=/var/lib/kubernetes/service-account-key.pem \\\
  --service-cluster-ip-range=10.32.0.0/24 \\\
  --use-service-account-credentials=true \\\
  --v=2
Restart=on-failure
RestartSec=5
```



# View controller-manager options

```
ps -aux | grep kube-controller-manager
```

```
root      1994    2.7   5.1 154360 105024 ?        Ssl  06:45   1:25 kube-controller-manager --  
address=127.0.0.1 --cluster-signing-cert-file=/etc/kubernetes/pki/ca.crt --cluster-signing-key-  
file=/etc/kubernetes/pki/ca.key --controllers=*,bootstrapsigner,tokencleaner --  
kubeconfig=/etc/kubernetes/controller-manager.conf --leader-elect=true --root-ca-  
file=/etc/kubernetes/pki/ca.crt --service-account-private-key-file=/etc/kubernetes/pki/sa.key  
--use-service-account-credentials=true
```

# Course Objectives

Core Concepts

Cluster Architecture

Services & Other Network Primitives

API Primitives

Scheduling

Logging Monitoring

Application Lifecycle Management

Cluster Maintenance

Security

Storage

Networking

Installation, Configuration & Validation

Troubleshooting

# Kube Scheduler





## Master

Manage, Plan, Schedule, Monitor  
Nodes



## Worker Nodes



Host Application as Containers

Controller-  
Manager


ETCD  
CLUSTER



Kube-Scheduler




# Kube-Scheduler



# Kube-Scheduler




# Kube-Scheduler



1. Filter Nodes

2. Rank Nodes



# View kube-scheduler options

## - kubeadm

```
▶ cat /etc/kubernetes/manifests/kube-scheduler.yaml
```

```
spec:
  containers:
  - command:
    - kube-scheduler
    - --address=127.0.0.1
    - --kubeconfig=/etc/kubernetes/scheduler.conf
    - --leader-elect=true
```



# View kube-scheduler options

```
▶ ps -aux | grep kube-scheduler  
root      2477  0.8  1.6 48524 34044 ?        Ssl  17:31   0:08 kube-scheduler --  
address=127.0.0.1 --kubeconfig=/etc/kubernetes/scheduler.conf --leader-elect=true
```

# Kubelet



kube-apiserver



## Master

Manage, Plan, Schedule, Monitor  
Nodes

kubelet



## Worker Nodes

Host Application as Containers

Controller-  
Manager

D  
CLUSTER

e-scheduler



# Kubernetes Architecture



## Master

Manage, Plan, Schedule, Monitor  
Nodes



## Worker Nodes

Host Application as Containers

kube-apiserver

ETCD  
CLUSTER

Controller  
-Manager

kube-scheduler

kubelet



kubelet



# Installing kubelet

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kubelet
```

## kubelet.service

```
ExecStart=/usr/local/bin/kubelet \\  
--config=/var/lib/kubelet/kubelet-config.yaml \\  
--container-runtime=remote \\  
--container-runtime-endpoint=unix:///var/run/containerd/containerd.sock \\  
--image-pull-progress-deadline=2m \\  
--kubeconfig=/var/lib/kubelet/kubeconfig \\  
--network-plugin=cni \\  
--register-node=true \\  
--v=2
```




# View kubelet options

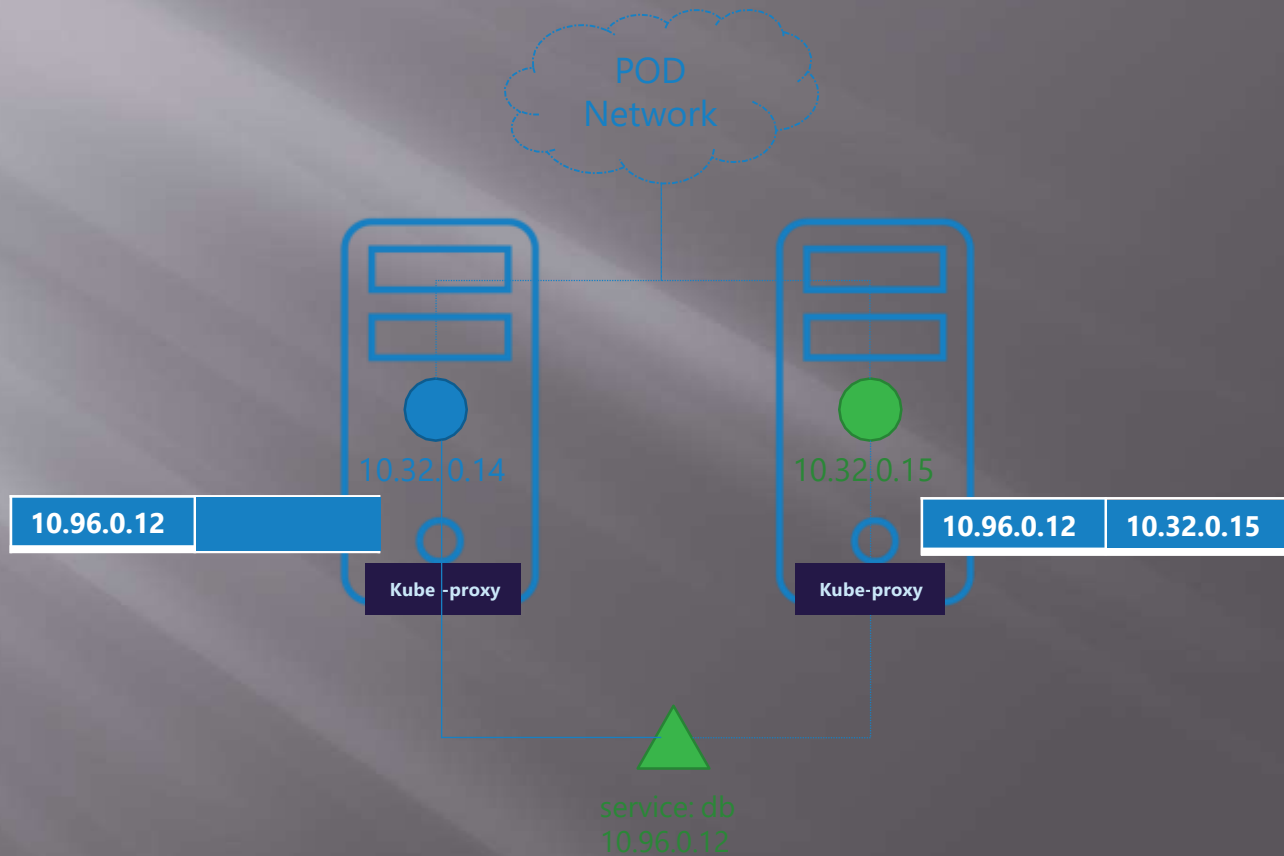
```
ps -aux | grep kubelet
```

```
root  2095  1.8  2.4  960676  98788  ?  Ssl  02:32  0:36  /usr/bin/kubelet  --bootstrap-  
kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --  
config=/var/lib/kubelet/config.yaml --cgroup-driver=cgroupfs --cni-bin-dir=/opt/cni/bin --cni-  
conf-dir=/etc/cni/net.d --network-plugin=cni
```

# Kube-proxy

An abstract network diagram in the bottom-left corner, featuring several white nodes connected by thin white lines, forming a web-like structure.

# Kube-proxy





# Installing kube- proxy

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kube-proxy
```

kube-proxy.service

```
ExecStart=/usr/local/bin/kube-proxy \\  
  --config=/var/lib/kube-proxy/kube-proxy-config.yaml  
Restart=on-failure  
RestartSec=5
```

# View kube-proxy - kubeadm

```
kubectl get pods -n kube-system
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcd6894-hwrq9	1/1	Running	0	16m
kube-system	coredns-78fcd6894-rzhjr	1/1	Running	0	16m
kube-system	etcd-master	1/1	Running	0	15m
kube-system	kube-apiserver-master	1/1	Running	0	15m
kube-system	kube-controller-manager-master	1/1	Running	0	15m
kube-system	kube-proxy-lzt6f	1/1	Running	0	16m
kube-system	kube-proxy-zm5qd	1/1	Running	0	16m
kube-system	kube-scheduler-master	1/1	Running	0	15m
kube-system	weave-net-29z42	2/2	Running	1	16m
kube-system	weave-net-snm1	2/2	Running	1	16m

```
kubectl get daemonset -n kube-system
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
kube-proxy	2	2	2	2	2	beta.kubernetes.io/arch=amd64	1h

NAME	NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
coredns-78fcd6894-hwrq9	kube-system	1/1	Running	0	16m	kube-system
coredns-78fcd6894-rzhjr	kube-system	1/1	Running	0	15m	kube-system
etcd-master	kube-system	1/1	Running	0	15m	kube-system
kube-apiserver-master	kube-system	1/1	Running	0	15m	kube-system
kube-controller-manager-master	kube-system	1/1	Running	0	15m	kube-system
kube-proxy-lzt6f	kube-system	1/1	Running	0	16m	kube-system
kube-proxy-zm5qd	kube-system	1/1	Running	0	16m	kube-system
kube-scheduler-master	kube-system	1/1	Running	0	15m	kube-system
weave-net-29z42	kube-system	2/2	Running	1	16m	kube-system
weave-net-snm1	kube-system	2/2	Running	1	16m	kube-system

AKS

	NOT CHARGED	CHARGED	OPTIONAL / ADDITIONAL
KUBERNETES CONTROL-PLANE	✓		
BILLING SUPPORT	✓		
BASIC LOAD BALANCER	✓		
NODE'S VIRTUAL MACHINES		✓	
NODE'S DISKS		✓	
BANDWIDTH		✓	
STORAGE FOR PERSISTENT VOLUMES			✓
STANDARD LOAD BALANCER			✓
PUBLIC IP ADDRESSES			✓

- ▣ “There are two ways to be successful with kubernetes:
- ▣ One: you’re goon be running certified kubernetes distributions everywhere.
- ▣ Two: You need to be able to manage all these clusters in a simple, consistent way.”

▣ -PETER SMAILES, CHIEF MARKETING OFFICER AT RANCER

# Benefits of Managed Kubernetes

- ❑ Multi-master cluster, including redundant etcd servers;
- ❑ Abstracts complexity of distributed systems;
- ❑ Makes things easier to run Kubernetes at scale;
- ❑ SLA and Support;
- ❑ Managed kubernetes allows you to focus on end-to-end DevOps and SRE solutions;

# Limitation

- ▣ Kubernetes Version;
- ▣ Predefined Container runtime and Container Network Interface (CNI).
- ▣ Predefined OS base image;
- ▣ Control Plane limited access and customization;
- ▣ Not all features are available in all regions;
- ▣ Service Limits;

# AKS Capacity Analysis

- ▣ Regions availability
- ▣ References
- ▣ Products available by region
- ▣ <https://azure.microsoft.com/en-us/global-infrastructure/services/?products=kubernetes-service>
- ▣ Restricted Virtual Machine Sizes
- ▣ References
- ▣ Restricted VM sizes
- ▣ <https://docs.microsoft.com/en-us/azure/aks/quotas-skus-regions#restricted-vm-sizes>
- ▣ Resource Group and Tags
- ▣ References
- ▣ Can I modify tags and other properties of the AKS resources in the node resource group?
- ▣ <https://docs.microsoft.com/en-us/azure/aks/faq#can-i-modify-tags-and-other-properties-of-the-aks-resources-in-the-node-resource-group>
- ▣ Node Operating System
- ▣ References
- ▣ Change kernel parameters
- ▣ <https://www.devops.buzz/public/aks/reboot-node#change-kernel-parameters>
- ▣ Kubernetes version
- ▣ References
- ▣ Supported Kubernetes versions in Azure Kubernetes Service (AKS)
- ▣ <https://docs.microsoft.com/en-us/azure/aks/supported-kubernetes-versions>



# AKS Capacity Analysis

- ▣ Network
- ▣ References
- ▣ Compare network models
- ▣ <https://docs.microsoft.com/en-us/azure/aks/concepts-network#compare-network-models>
- ▣ Deep dive – Azure Load Balancer
- ▣ <https://msandbu.org/deep-dive-azure-load-balancer/>
- ▣ Resources reservation
- ▣ References
- ▣ Resource reservations
- ▣ <https://docs.microsoft.com/en-us/azure/aks/concepts-clusters-workloads#resource-reservations>
- ▣ Azure subscription limits
- ▣ References
- ▣ Classic deployment model limits
- ▣ <https://docs.microsoft.com/en-us/azure/azure-resource-manager/management/azure-subscription-service-limits#classic-deployment-model-limits>
- ▣ Resource Group limits
- ▣ References
- ▣ Resource group limits
- ▣ <https://docs.microsoft.com/en-us/azure/azure-resource-manager/management/azure-subscription-service-limits#resource-group-limits>
- ▣ Resources not limited to 800 instances per resource group
- ▣ <https://docs.microsoft.com/en-us/azure/azure-resource-manager/management/resources-without-resource-group-limit>

# AKS network architecture

- ▣ References
- ▣ Network concepts for applications in Azure Kubernetes Service (AKS)  
<https://docs.microsoft.com/en-us/azure/aks/concepts-network>
- ▣ Create a private Azure Kubernetes Service cluster  
<https://github.com/MicrosoftDocs/azure-docs/blob/master/articles/aks/private-clusters.md>
- ▣ Kubenet (basic) vs Azure CNI (advanced)
- ▣ References
- ▣ Compare network models  
<https://docs.microsoft.com/en-us/azure/aks/concepts-network#compare-network-models>
- ▣ Basic Load Balancer vs Standard Load Balancer
- ▣ References
- ▣ Deep dive – Azure Load Balancer  
<https://msandbu.org/deep-dive-azure-load-balancer/>